

This document provides a detailed breakdown of the `index.html` file, explaining how its HTML, CSS, and JavaScript work together to create a functional chat interface that meets all your assignment requirements.

## 1. HTML Structure (The Skeleton)

The HTML provides the basic layout and elements of the chat application.

- `<head>` :
  - Imports the "Inter" font from Google Fonts to match the modern UI style.
  - Contains all the CSS within a `<style>` tag, as requested instead of SCSS.
- `<body>` :
  - Uses Flexbox (`display: flex`) to center the main chat window vertically and horizontally on the page.
- `.chat-container` :
  - This is the main white box that holds the entire chat UI. It uses Flexbox in a `column` direction to stack the header, message list, and input box on top of each other.
- `.chat-header` :
  - The top bar containing the title (`<h3>ZUS AI Agent</h3>`) and the `/reset` button.
- `.message-list` :
  - The main, scrollable area where messages appear.
  - `flex-grow: 1` is the key property that makes this element expand to fill all available space, pushing the composer to the bottom.
  - `overflow-y: auto` makes it scrollable only if the messages overflow.
- `.composer-container` :
  - A wrapper that holds both the chat input and the (initially hidden) autocomplete box.

- `position: relative` is crucial for allowing the autocomplete box to be positioned above it.
- `.autocomplete-box :`
  - `Positioned absolute` to float on top of the UI, just above the input.
  - `display: none` hides it by default.
- `.chat-composer (the <form> ):`
  - The form at the bottom containing the `textarea` and the "Send" button .
- `<textarea id="message-input">` :
  - The multiline input field. `rows="1"` sets its initial height to one line, and it grows automatically using JavaScript.
- **SVG Icon ( `<svg id="icon-tool">` )**
  - An SVG "tool" icon is defined at the bottom. This is used in the JavaScript to visually represent agent steps without needing an external image file.

## 2. CSS Styling (The Appearance)

The CSS inside the `<style>` tag styles the HTML to look like your design.

- `:root` : A best-practice way to define CSS variables (like a theme). This makes it easy to change colors (e.g., `--brand-color`) in one place.
- `.message :`
  - This is the base class for a chat bubble wrapper. It uses `display: flex` to put the avatar and the message content side-by-side.
- `.message.user & .message.bot :`
  - These "modifier" classes change the styling.
  - `.message.user` uses `align-self: flex-end` to move the message to the **right**.
  - `.message.bot` uses `align-self: flex-start` to keep it on the **left**.

- They also apply different background colors and `border-radius` to create the "chat bubble" tail effect.
- `.agent-steps & .step` :
  - This is the visualization for your agent's plan.
  - `.agent-steps` is a container with a `border-left` that creates the vertical line, connecting the steps.
  - `.step` is the box for a single tool call, using flexbox to align the icon, name, and (most importantly) the `<pre>` tag.
  - `<pre> tag`: This is used to display the tool's JSON arguments. It's styled to look like a code block and uses `white-space: pre-wrap` and `word-break: break-all` to ensure long JSON strings wrap neatly instead of breaking the UI.
- `.composer-container & .chat-composer` :
  - The `<textarea>` is styled to auto-resize. Its `max-height` is set to 6 lines to prevent it from growing and taking over the whole screen.
- `.autocomplete-box` :
  - Styled to appear as a floating box above the composer.
  - The `.selected` class, which is managed by JavaScript, adds a highlight to the currently selected autocomplete item.

### 3. JavaScript Logic (The Brain)

This is the most complex part. It's wrapped in a `DOMContentLoaded` listener to ensure it only runs after all the HTML is loaded.

#### Core State & Config

- `messages = []` : An array that holds the entire chat history as JavaScript objects. This is your "source of truth."
- `SESSION_ID` : A unique ID generated on page load, sent to the backend with each request.
- `BACKEND_URL` : The URL for your FastAPI server.
- `QUICK_ACTIONS` : An array of objects defining your / slash commands.

## Key Functions

1. `loadChatHistory()` & `saveChatHistory()` (**Requirement Fulfilled**)
  - `loadChatHistory()` is called on page load. It checks `localStorage` for "chatHistory".
  - If history exists, it parses the JSON and calls `renderMessage()` for each saved message.
  - If no history exists, it creates a new welcome message.
  - `saveChatHistory()` is called after *any* change to the `messages` array (adding a user message, receiving a bot response). It uses `JSON.stringify()` to save the array to `localStorage`.
2. `addMessage(content, role)` & `renderMessage(message)`
  - `addMessage` is a helper function. It creates a new message object (with an ID, timestamp, etc.), adds it to the `messages` array, and then calls `renderMessage`.
  - `renderMessage` is the workhorse. It creates the HTML elements for a message, including the avatar, author, and timestamp.
  - It correctly handles rendering messages from history, including recreating the `steps` and `finalAnswer` from the saved data.
  - It creates an empty `<p>` tag with a loading spinner for new bot messages, which will be filled in by the streaming function.
3. `handleSubmit(e)`
  - Triggered by pressing "Enter" or clicking the Send button.
  - It gets the text, adds the user's message (e.g., `addMessage("Hi", "user")`), and then saves the history.
  - It clears the input text.
  - It creates the empty "bot" message shell (e.g., `addMessage("", "bot")`).
  - Finally, it calls `streamAgentResponse()`, passing in the user's text and the bot message data.
4. `streamAgentResponse(text, botMessageData)` (**The Core Logic**)

- This function uses the modern `fetch` API to make the `POST` request to your backend.
- It gets a `ReadableStream` from the `response.body` and uses a `TextDecoder` to read the incoming data in chunks.
- **Stream Parsing:** This is the most complex part. Your backend sends a stream of `<step>...</step>` tags.
  - It uses a `buffer` to store incoming text.
  - It repeatedly checks the buffer for a *complete* pair of `<step>` and `</step>` tags.
  - When it finds a complete step, it parses the content *inside* that step using string matching (`.match(/<step_name>(.*)</step_name>/)`) to find the tool name and the JSON arguments.
  - **Visualizing Steps:** If the `stepName` is **not** `final_answer`, it calls `createStepElement()` to build the visual step and appends it to the bot's message. It also saves this step to the `botMessageData.steps` array.
  - **Final Answer:** If the `stepName` is `final_answer`, it parses the JSON, extracts the `answer`, and puts that text into the bot's final `<p>` tag.
- **Error Handling:** A `try...catch` block wraps the entire fetch request. If the API fails (like your HTTP 500 test), it will catch the error and display a graceful error message in the chat.
- **Saving:** After the stream is complete (`done` is true), it calls `saveChatHistory()` one last time to save the *complete* bot response (with all its steps and the final answer).

## 5. Autocomplete & Input Handlers (Requirement Fulfilled)

- `handleInput(e)` : Fires on every keystroke.
  - It auto-resizes the textarea's height based on its content (`scrollHeight`).
  - If the text starts with `/`, it filters the `QUICK_ACTIONS` list and calls `showAutocomplete()`.
- `handleKeydown(e)` :
  - Handles "Enter-to-send" and "Shift+Enter-for-newline".
  - Handles "ArrowUp" and "ArrowDown" to navigate the autocomplete menu by changing `autocompleteIndex` and updating the `.selected` class.

- Handles "Enter" to select an item from the autocomplete menu.
- handleReset() : Clears the messages array, clears localStorage , and clears the UI.